

# Cloud4Things

Marcus Vinícius Paulino Gomes

Instituto Superior Técnico, Universidade de Lisboa  
`marcus.paulino.gomes@tecnico.ulisboa.pt`



## Table of Contents

1	Introduction.....	7
1.1	Internet of Things .....	8
1.2	Cloud Computing .....	9
2	Objectives .....	12
3	Related Work .....	15
3.1	Cloud Computing and Internet of Things.....	15
3.2	Automated Deployment of Cloud Applications .....	17
3.3	Service Level Agreements in IoT Applications .....	18
4	Solution Architecture .....	22
4.1	State of the Art .....	23
4.2	Cloud of Things Architecture .....	23
5	Evaluation Methodology .....	27
5.1	Cloud Performance Evaluation .....	27
5.2	RFID Evaluation .....	27
6	Conclusion .....	30
A	Planning .....	32



**Abstract.** In this document we present Cloud4Things, a solution that proposes to make the deploy and manage IoT applications for smart places. Due to the heterogeneity of IoT applications, management tasks such as application deployment and management are complex tasks that require advanced technical knowledge. Cloud4Things proposes to decrease the complexity of these tasks by adopting a high-level perspective enables to perform the deployment and management operations of IoT applications. To simplify the deployment operation at the Cloud, Cloud4Things will rely on cloud orchestrators tools to execute such task. These orchestration tools allow the specification of the application structure in high-level abstraction and allow the control the life-cycle of the application. Cloud4Things will allow the management of the provisioned resources at the Cloud by a smart place only having in mind the business rules of this particular space, thus smart places managers will be able to monitor the performance of its smart place in real time. This document also presents the state-of-the-art solution as well an overview about the proposed approaches and evaluation methods.

**Keywords:** Internet of Things, Cloud Computing, Smart Places, Service Level Agreements



## 1 Introduction

In recent years, computing is becoming more ubiquitous in the physical world. This notion of ubiquitous computing was introduced by Weiser [1], where the concept of smart environment was defined as “*a physical world richly and invisibly interwoven with sensors, actuators, displays, and computational elements, embedded seamlessly in the everyday objects of our lives and connected through a continuous network*”. Through the years, technology advances such as the creation of the Internet contributes to achieving the ubicomp’s view which enables individual devices to communicate between themselves from any part of the world [2]. In Weiser vision of ubiquitous computing is a calm world where information seamlessly move in and out of attention as automation gives way to human interaction [3].

However, due to the slow progress in concretize this vision [4], another approaches for ubiquitous computing that constrasts with Weiser vision becomes relevant. Rogers proposes an approach which focuses in designing technologies for engaging user experiences that in a creative and constructive way extend the peoples capabilities [5]. Rogers points that ubiquitous technologies can be developed not only for individuals, but for particular domains that can be set up and customized by an organization according its needs, such as agriculture, retailing and logistics.

Actually, Weiser vision is close to becoming reality thanks to the Internet of Things and Cloud Computing [6]. This world where things are connected through a continuous network is a vision thats represents the Internet of Things (IoT). In this vision, physical items are continuously connected to the virtual world and can act as remotely physical access points to Internet Services. The Internet of Things make computing truly ubiquitous [7].

A common scenario where the Internet of Things paradigm is applied are smart environments [8], which in this document are designated as smart places. In the context of this work, a smart place can be defined as an ecosystem composed by smart objects such as RFID tags and sensors, that are able to acquire knowledge about this environment and also to adapt this inhabitants in order to improve their experience in that environment [9]. In particular, the deployment of IoT applications in a smart place is a challenge today due the heterogeneity of the smart objects that are present in a smart place and also because the required infrastructure. In order to decrease the complexity of the deployment of an IoT application in a smart place, the Cloud computing paradigm allows to virtualize the required physical infrastructure. Nowadays the infrastructure needed by IoT applications can be virtualized by Cloud providers, which helps to gain more flexibility and reduce the costs in the deployment of an IoT application. Therefore, another challenge concerns with the heterogeneity of a smart place, more precisely the variety of smart objects that can be inside of a given place. Many of these objects uses different communication protocols and drivers to commu-

nicate with devices, thus the configuration of these objects must be manually handled, which makes the integration of these objects in the deployment of IoT applications in a smart place an inefficient process.

In order to create a more solid background about the concepts that are around of a smart place, the next sections are dedicated to give a more detailed description of these concepts such as Internet of Things and Cloud Computing.

### 1.1 Internet of Things

The Internet of Things is a concept in which the virtual world of information technology integrates seamlessly with the real world of things [10]. Through the amount of computer and network devices available nowadays, the real world becomes accessible to business and everyday scenarios. A more precise definition of what does the Internet of Things was formulated in the Strategic Research Agenda of the Cluster of European Research Projects on the Internet of Things (CERP-IoT 2009):

*“In the IoT, ‘things’ are expected to become active participants in business, information and social processes where they are enabled (...) by exchanging data and information sensed about the environment, while reacting autonomously to the ‘real/physical world’ events and influencing it by running processes that trigger actions and create services with or without direct human intervention. Interfaces in the form of services facilitate interactions with these ‘smart things’ over the Internet (...) taking into account security and privacy issues.”*

IoT provides access to a more detailed information, which results that the level of analysis can be performed in a large-scale perspective, as well in a small-scale perspective. But Internet of Things is more than a tool for managing business processes more efficiently and more effectively, IoT will also enable a more convenient life for all peoples. Recently, IoT became relevant to industry and end-users [10], mainly because the reduction of cost and miniaturization of technologies used in IoT applications, such as RFID, sensor networks, NFC and wireless communication. Technologies like these allow the detection of status of things, which together with collection and processing of detailed data, allows to create an interactive and responsive network with huge potential for citizens, consumers, business and where is possible to gather immediate responses to events that occurs in the real world [10]. One of the expectations regarding of IoT concerns with real-world awareness provided to information systems [7]. An example is in logistics applications, by the use of RFID, companies can react promptly to relevant physical events, and then manage their processes in a better way, typically increasing efficiency and reducing costs. Another expectation of IoT concerns with providing services to the end-users through common objects, in that way products will be able to provide recommendations for use and maintenance instructions, supply warranty information or highlight complementary products.



## 1.2 Cloud Computing

Cloud Computing is a paradigm where its infrastructure is designated as a “Cloud” from which its resources are available to businesses and users from anywhere in the world on demand [11]. This paradigm enables rapid service delivery in a dynamically scalable and virtualized manner. Actually, clouds are built on top of modern data centers, that incorporated Infrastructure as a Service (IaaS), Platform as a Service (PaaS) and Software as a Service (SaaS) [12], as illustrated on Figure 1. Cloud providers, such as Amazon Web Services<sup>1</sup>, Google Cloud Platform<sup>2</sup> and Microsoft Azure<sup>3</sup> offers these services as utilities in a pay-per-use model in which Quality of Service *QoS* guarantees are offered by means of customized Service Level Agreements (SLA’s) [13], that are negotiated between the customers and the Cloud providers.




Fig. 1: Hierarchical View of Cloud Computing.

Data Centers allocate all the hardware that supports the cloud, they consist of thousands of physical machines that provide redundancy and ensure reliability in case of site failures. By using the IaaS pattern, Cloud Computing is able to virtualize resources as computing power, storage and network connectivity of the data centers. These computational resources are provisioned on demand in a form of Virtual Machines (VM’s) deployed in a cloud provider data-center [14]. Cloud also assists application design, development, application hosting and testing by providing a development platform in order to assist the execution of these tasks, this development as a service pattern is known as PaaS. Another pattern supported by Cloud Computing is SaaS. This pattern allows that a single piece of software is transferred to millions of users through a browser [15]. This pattern is beneficial by both users and developers, the developers only need to

<sup>1</sup> <http://www.aws.amazon.com>

<sup>2</sup> <http://cloud.google.com>

<sup>3</sup> <http://www.azure.microsoft.com/>

maintain a single program while the users can save costs in servers and software.

In a global analysis Cloud computing offers several benefits such as high-availability, high-scalability, flexibility, on-demand service provisioning and massive reduction of costs. By converging IoT applications and Cloud computing, IoT applications are able to take advantage of all these benefits.

**Cloud Orchestration** Due the heterogeneity of IoT applications infrastructure service management tasks such as deployment, driver installation and gateway configuration are still handled manually in a particular way for each case. In order to reduce the complexity of this tasks, Cloud Orchestration tools can be used to automate these management tasks. The process of orchestration woven the software components of the application in a one piece that can be managed more effectively, in order to ensure a smooth and fast service delivery. Cloud Orchestration helps to take advantage of the full benefits of Cloud computing by providing features that can include:

- Simplify, automate and optimize service deployment by integrating cloud capabilities across heterogeneous environments.
- Automated high-scale provisioning and de-provisioning of resources with policy-based tools.
- Real-time monitoring of physical and virtual Cloud resources.
- Selection of Cloud services such as storage and hosting, through a self-service portal.
- Real-time monitoring of usage and accounting chargeback capabilities to track and optimize system usage.

In addition, orchestration allows to reduce significantly the costs related with labor and resources, since that manual intervention and management of varied IT services and resources are not needed. Actually several tools to performing Cloud Orchestration are available in the market such as IBM Cloud Orchestrator<sup>4</sup>, HP Operations Orchestration<sup>5</sup> and GigasSpaces Cloudify<sup>6</sup> and some open-source tools like Ubuntu Juju<sup>7</sup> and OpenTOSCA<sup>8</sup>.

The remainder of this document is organized as the follows. In Section 2 we describe the problem to be solved as well the main objectives of this work. Section 3 presents the related work in the research area. Then in Section 4 we present a brief description of the state of art solution and propose the architecture for our solution. In Section 5 the evaluation methodology is described. In the Section 6 we presents the conclusion and finally in the appendix we present the schedule for the future work.

<sup>4</sup> <http://www.ibm.com/software/products/en/ibm-cloud-orchestrator>

<sup>5</sup> <http://www8.hp.com/us/en/software-solutions/operations-orchestration-it-process-automation>

<sup>6</sup> <http://www.gigaspaces.com/cloudify-cloud-orchestration/overview>

<sup>7</sup> <http://www.juju.ubuntu.com>

<sup>8</sup> <http://www.iaas.uni-stuttgart.de/OpenTOSCA/indexE.php>



## 2 Objectives

The life-cycle of a Cloud-based IoT application is composed by several stages. As illustrated in the Figure 2, some of the stages are placed in the smart place while others are placed in the Cloud.




Fig. 2: IoT application life-cycle.

The life-cycle of an IoT application starts with the installation of the sensors and readers in the smart place (step 1). The next step consists in prepare the smart place in order to process the events sent by the smart objects. First the application must be deployed at the Cloud in order to be able to receive the events (step 2). After that the installed sensors and readers must be configured to sent the generated events in the smart place to the application that is running in the Cloud (step 3). At this point the smart place is already configured to support the processing of events generated by smart objects (step 4). A very important point is to assure that the smart place is working according with the desired Quality of Service (QoS). QoS is a concept that embraces a number of nonfunctional properties such as price, availability, reliability and reputation [16]. Thus in order to monitoring the performance of the application in the smart place (step 5), these properties must be agreed between the customer and the service provider, in this case the Cloud provider. This agreement between the customer and the service provider is know as Service Level Agreement (SLA). The SLA defines the terms and conditions of service quality that a service provider delivers to service requesters based on the *QoS* information [17]. The final stage that an IoT application can reach is its undeployment, which means that the smart

place is deactivated (step 6).

Thus, the objective of this work is to decrease the complexity of deployment and management of Cloud-based IoT applications in a smart place. Usually the deployment of such applications is performed by a technician that manually configure and install the components of the application. In order to reduce the complexity of this process, the most effective approach is to automate the deployment process of the application. That is achieved by performing the deployment through Cloud Orchestration tools, that allows to specify the components and the relations between themselves in a high-level perspective and also provisioning the necessary resources at the Cloud in a effective way. These tools also allows to perform the monitoring of these applications during its life-cycle as well undeploying them. However, this tools don't solve all our problems. In particular, IoT applications require a software stack that usually is composed by a database, a web server and the event processing software. But orchestration tools normally lacks the integration of the event processing software with the other components. Thus, to support the integration of such components these tools must be extend to support the event processing software.

However, managing these applications through a Orchestrator tool requires an elevated technical knowledge. To allow that non-technical users can perform the managing operations having in perspective the high-level business rules of the smart place, the main objective of this work is to permit that non-technical users can perform the management of smart places only having in mind the business rules of that particular place. To achieve that, these high-level rules must be translated to a more low-level rules that can be expressed in terms of non-functional requirements, SLAs, that can be used to estimate the resources needed by the application in order to have an acceptable *QoS*. In the other hand, to allow those non-technical users to monitoring the smart place, these SLAs must be translated to a high-level rules that can be expressed in terms of business rules of a given smart place. For instance, if a smart place has a flow of people of 500 persons per day, that business rule will be translated to a SLA that express the amount of storage and bandwidth required to ensure the *QoS* of the application. Furthermore, by monitoring the service level offered by the Cloud providers it will be possible to determine if the Cloud is overloaded with the amount of data generated by the smart place or not.



### 3 Related Work

Cloud and Internet of Things are emerging computing paradigms that features distinctively different computing resources and system architecture. As its popularity has been growing across the academics and the industry, researchers and developers are spending a lot of effort to investigate how to integrate these technologies in order to take advantage of the benefits provided by both of them.

IoT applications often encapsulate several relatively complex protocols and involves different software components. Moreover, they require a significant investment in infrastructure, besides that the system administrators and users spend time with large client and server installations, setups, or software updates. As most of the computing resources are allocated on the Internet on servers in Cloud computing, integrating these paradigms in a Cloud-based model results in a solution with more flexibility of implementation, high scalability and high availability, and with a reduced upfront investment.

**Fosstrak** The capture of events will be performed by Fosstrak. Fosstrak is an open-source RFID project<sup>9</sup> that implements the EPC (Electronic Product Code) Network specifications. Fosstrak is composed by modules that allows to store and query about events captured by a RFID reader, convert the EPC identifiers to others EPC representations, filter and collect the data from RFID readers and to configure and manage RFID readers. These modules provides an entire RFID infrastructure that accelerates the development of IoT applications and allows to users to gain hands-on experience with the EPC Network.

#### 3.1 Cloud Computing and Internet of Things

In RFID-based IoT applications, Guinard et al. [18] point out that the deployment of RFID applications are cost-intensive mostly because they involve the deployment of often rather large and heterogeneous distributed systems. As a consequence, these systems are often only suitable for big corporations and large implementations and do not fit the limited resources of small to mid-size businesses and small scale applications both in terms of required skill-set and costs. To address this problem, Guinard et al. proposes a Cloud-based solution that integrates virtualization technologies and the architecture of the Web and its services. The case of study consists in a IoT application that uses RFID technology to substitute existing Electronic Article Surveillance (EAS) technology, such as those used in clothing stores to track the products. In this scenario they applied the Utility Computing blueprint to the software stack required by the application using the AWS platform and the EC2 service. The Elastic Cloud Computing (EC2) service allows the creation and management of virtual machines (Amazon Machine Images, or AMIs) that can then be deployed on demand onto a

---

<sup>9</sup> <https://code.google.com/p/fosstrak/>

pool of machines hosted, managed and configured by Amazon. A benefit of this approach is that the server-side hardware maintenance is delegated to the cloud provider which is often more cost-efficient for smaller businesses. Furthermore, it also offers better scaling capabilities as the company using the EPC Cloud AMI, can deploy additional and more powerful instances regarding to the amount of requests.

Distefano [19] et al. proposed a high-level modular architecture to implement the Cloud of Things. According to Distefano et al. things not only can be discovered and aggregated, but also provided as a service, dynamically, applying the Cloud provisioning model to satisfy the agreed user requirements and therefore establishing Things as a Service providers. The *Things as a Service* (TaaS) paradigm envisages new scenarios and innovative, pervasive, value-added applications, disclosing the Cloud of Things world to customers and providers as well, thus enabling an open marketplace of "things". To address this issues, an ad-hoc infrastructure is required to deal with the management of sensing and actuation, mashed up resources provided by heterogeneous Clouds, and things, by exploiting well known ontologies and semantic approaches shared by and adopted by users, customers and providers to detect, identify, map and transform mashed up resources. The proposed architecture provides blocks to deal with all the related issues, while aiming to provide things according to a service oriented paradigm.

CloudThings [20] is an architecture that uses a common approach to integrate Internet of Things and Cloud Computing. The proposed architecture is an online platform which accommodates IaaS, PaaS, and SaaS and allows system integrators and solution providers to leverage a complete IoT application infrastructure for developing, operating and composing IoT applications and services. The applications consists of three major modules, the CloudThings service platform, that is a set of Cloud services (IaaS), allowing users to run any applications on Cloud hardware. This service platform dramatically simplifies the application development, eliminates need for infrastructure development, shortens time to market, and reduces management and maintenance costs. The CloudThings Developer Suite is a set of Cloud service tools (PaaS) for application development, such as Web service API's, which provide complete development and deployment capabilities to developers. The CloudThings Operating Portal is a set of Cloud services (SaaS) that support deployment and handle or support specialized processing services. To evaluate CloudThings, a smart home application based on a Cloud infrastructure was implemented. In the application, the sensors read the home temperature and luminosity and the Cloud application stores and visualized them, so that the user can view the smart home temperature and luminosity anywhere. In particular, in these implementation the Cloud architecture was extended by inserting a special layer for dynamic service composition. This middleware encapsulates sets of fundamental services for executing the users service requests and performing service composition, such as process planning, service



discovery, process generation, process execution, and monitoring. This first implementation also demonstrates that this middleware as a service releases the burden of costs and risks for users and providers in using and managing those components.

### 3.2 Automated Deployment of Cloud Applications

The effort put in the research to integrate the paradigms of Cloud Computing and Internet of Things resulted in a essential contribution, but there are several issues regarding to the integration between Cloud Computing and Internet of Things that must be addressed. In particular, due of the heterogeneity of the IoT applications environments, its hard for solution providers to efficiently deploy and configure applications for a large number of users. Thus, automation for the management tasks required by IoT applications is a key issue to be explored.

TOSCA (Topology and Orchestration Specification for Cloud Applications) [21] is proposed in order to improve the reusability of service management processes and automate IoT application deployment in heterogeneous environments. TOSCA is a new cloud standard to formally describe the internal topology of application components and the deployment process of IoT applications. The structure and management of IT services is specified by a meta-model, which consists of a *Topology Template*, that is responsible to describe the structure of a service, then there are the *Artifacts*, that describes the files, scripts and software components necessary to be deployed in order to run the application, and finally the *Plans*, that defined the management process of creating, deploying and terminating a service. The correct topology and management procedure can be inferred by a TOSCA environment just by interpreting te topology template, this is known as "declarative" approach. Plans realize an "imperative" approach that explicitly specifies how each management process should be done. The topology templates, plans and artifacts of an application are packaged in a Cloud Service Archive (.csar file) and deployed in a TOSCA environment, which is able to interpret the models and perform specified management operation. These .csar files are portable across different cloud providers, which is a great benefit in terms of deployment flexibility. To demonstrate its feasibility TOSCA was used to specify a typical IoT application in building automation, an Air Handling Unit (AHU). The common IoT components, such as gateways and drivers will be modeled, and the gateway-specific artifacts that are necessary for application deployment will also be specified. By archiving the previous specifications and corresponding artifacts into a .csar file, and deploying it in a TOSCA environment, the deployment of AHU application onto various gateways can be automated. As a newly established standard to counter growing complexity and isolation in cloud applications environments, TOSCA is gaining momentum in industrial adoption as well academic interests.

Breitenbücher et al. [22] proposed to combine the two flavors of management supported by TOSCA, *declarative processing* and *imperative processing*, in order to create a standards-based approach to generate provisioning plans based on TOSCA topology models. The combination of both flavors would enable applications developers to benefit from automatically provisioning logic based on declarative processing and individual customization opportunities provided by adapting imperative plans. These provisioning plans are workflows that can be executed fully automatically and may be customized by application developers after generation. The approach enables to benefit from strengths of both flavors that leads to economical advantages when developing applications with TOSCA. The motivating scenario that is used to evaluate this approach consists in a LAMP-based TOSCA application to be provisioned. The application implements a Web-shop in PHP that uses a MySQL database to store product and customer data. The application consists of two application stacks, one provides the infrastructure for the application logic and the other hosts the database, were both will be run on Amazon's public IaaS offering Amazon EC2. To measure the performance of the deployment using the two-flavor approach, the strategy adopted was measure the time spent to generate provisioning plans regarding the number of templates required by the application. The results indicates that the required time increases linearly to the number of templates.

Recently a growing number of organizations are developing Orchestrators, Design Tools and Cloud Managers based on TOSCA. Juju is an Open Source TOSCA Orchestrator that can deploy workloads across public, private clouds, and directly onto bare metal. HP Cloud Service Automation is cloud management solution that supports declaratives services design that are aligned with TOSCA modeling principles. GigaSpaces Cloudify orchestrates TOSCA Service Templates using workflows to automate deployments and other DevOps automation processes. IBM Cloud Orchestrator provides integrated tooling to create TOSCA applications, deploy them with custom policies, monitoring and scale them in cloud deployments.

### 3.3 Service Level Agreements in IoT Applications

Initially the objective of this work was focused in improve the deployment operation. However, during the research of this work it became clear that the automation of the deployment of IoT applications in smart places was an area that has a notorious progress in the developed work, standing out the development of Cloud Orchestration tools based on TOSCA. Therefore, it was realized that the focus of this work could be redirected to explore a field that still needs large research effort and is fundamental for the IoT success [8], *QoS* support. In order to guarantee an acceptable *QoS* for Cloud-based IoT applications, SLAs must be negotiated between service customers and services providers in order to guarantee an acceptable *QoS* for the service. Recently, *QoS* support has been a field where researchers are proposing new approaches regarding policy-aware IoT

applications based on TOSCA.

Waizenegger et al. [23] proposed a mechanism to demonstrate how non-functional requirements are defined in TOSCA using policies. TOSCA allows the specification of non-functional requirements like cost, security, and environmental issues. However, TOSCA lacks detailed description of how to apply, design, and implement policies. This paper proposes two mechanisms for processing policies during the deployment and management of Cloud services in TOSCA. The policies are described according to a taxonomy where a Cloud service policy is a tuple of elements that describe its behavior and effect. To perform the implementation of this policy-specific logic two approaches were proposed, an Implementation Artifact Based Policy Enforcement (IA-Approach) and a Plan-Based Policy Enforcement (P-Approach) [23]. In the IA-Approach the existing Implementation Artifacts - responsible to perform service management operations - are extended to support policy enforcement implementations, thus these artifacts are comprised of two alternative implementations for each service management operation: one that is policy enforcing and one that is not. In the P-Approach the policies are enforced in an imperative way, for each policy the plan executes the appropriate steps to enforcing the policy. Although both of these approaches allow the enforcement of policies, the approach used for realizing this should be chosen with caution. The IA-Approach is more suitable for less complex scenarios and the policies implemented with this approach can be reused very easily. On the other hand, the enforcement of policies in more complex scenarios should be performed according to the P-Approach, that allows to have a global control of the triggered enforcement actions.

Breitenbücher et al. [24] proposed an approach that enables to automate the provisioning and management of composite Cloud applications in compliance with non-functional security requirements defined by policies. In this approach the Management Planlet Framework [25] was extended to support policy-aware provisioning and management based on Management Policies that bind non-functional security requirements to the management tasks that must enforce them. This paper also introduces the concept of Management Annotation Policies, which define the semantics and how the policy must be processed. The automation of provisioning and management tasks are performed by Policy-aware Management Planlets, that execute these tasks considering the Annotation Policies attached to each operation. The Framework also allows the Cloud providers and application developers to specify their own policy-aware management logic in a flexible and reusable manner independently from individual applications. The management logic can be defined in two ways, by applying Automated Management Patterns or by manually creating Desired Application State Models. The Framework was evaluated in terms of feasibility, performance, economics, limitations and extensibility. To validate the concept technically, a prototype was successfully implemented based on a previous implementation of the Manage-

ment Planlet Framework.

#### Colocar LoHiFramework

An important observation of the realized work until now is that the developed approaches to enable policy-aware IoT applications are based in low-level policies. Thus, in the interest to add a new contribute to this field, the objective of this work is focused in develop an approach where these policy-aware applications are based in high-level business rules.



## 4 Solution Architecture

Actually a typical scenario for an IoT application consists of a smart space that contains several smart objects and all the infrastructure required to support this application, namely RIFD tags, sensors, readers and servers, as illustrated at the Figure 3. This scenario presents several issues regarding the deployment of the application, the low scalability, the costs of infrastructure and the maintenance of the same.

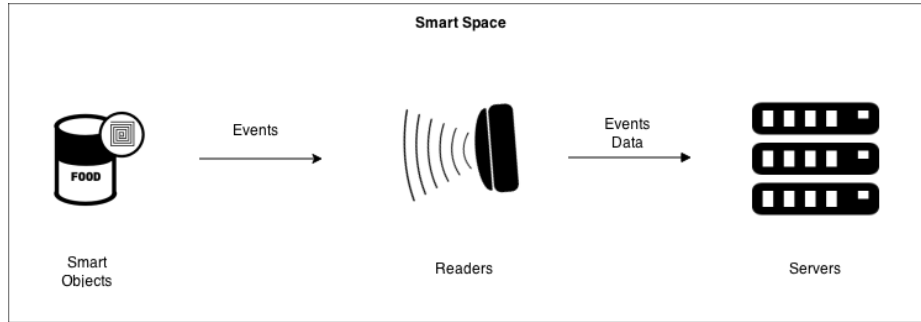


Fig. 3: Typical smart space scenario.

By converging the IoT applications with the Cloud Computing paradigm the objective is simplify this scenario by leveraging the required infrastructure by these applications to the Cloud providers, as illustrated at the Figure 4. Furthermore, the convergence of this two paradigms, allows to take advantage of the benefits offered by Cloud computing as referenced in **Section 1.2**.




Fig. 4: Cloud-based smart space.

#### 4.1 State of the Art

Actually Cloud-based IoT applications are the state-of-the-art of this kind of solutions. By leveraging the infrastructure to the Cloud providers, these applications has a high-availability, can dynamically scale while spending a fraction compared with the traditional solutions. However, as earlier mentioned the deployment of such applications still is an issue, due its complexity and required manual intervention. The deployment of Cloud-based IoT applications usually are performed through IT automation tools, such as Chef<sup>10</sup> and Puppet<sup>11</sup>. These tools enables to automate the deployment of such applications in a certain way, given that all the components of the application and the relation between themselves must be specified manually, which requires considerable manual work and expertise by the person that is performing the deployment. However, the deployment process of these solutions is not the only existent issue. Monitoring the application life-cycle is a task that requires a lot of effort and expertise by the system administrators.

In order to solve this problem, the adopted approach relies on perform the deployment of these applications by using Cloud Orchestrator tools. As mentioned in **Subsection 1.2**, these tools allows to specify the application components and their relations in a high-level perspective and to execute the management tasks required by the application during its life-cycle. As a matter of fact, in a low-level perspective cloud orchestration tools express the high-level perspective defined by the user into scripts that latter are executed using IT automation tools such as Puppet and Chef.

#### 4.2 Cloud of Things Architecture

The main objective of Cloud of Things is decrease the complexity of deployment and management of IoT applications. To achieve this objective, Cloud of Things must enable non-technical users - the business managers - to perform the monitoring of Cloud-based IoT applications as well defining Service Level Agreements in a high-level way. In order execute this tasks, users must be able to interact with the application that is running at the cloud in order to observe their state and to apply some decisions based on the performance of the application. Thus, Cloud of Things must provide a service, e.g a GUI, that allows the users to perform such actions. At Figure 5 we present the Cloud of Things architecture.

---

<sup>10</sup> <http://www.chef.io>

<sup>11</sup> <http://www.puppetlabs.com>

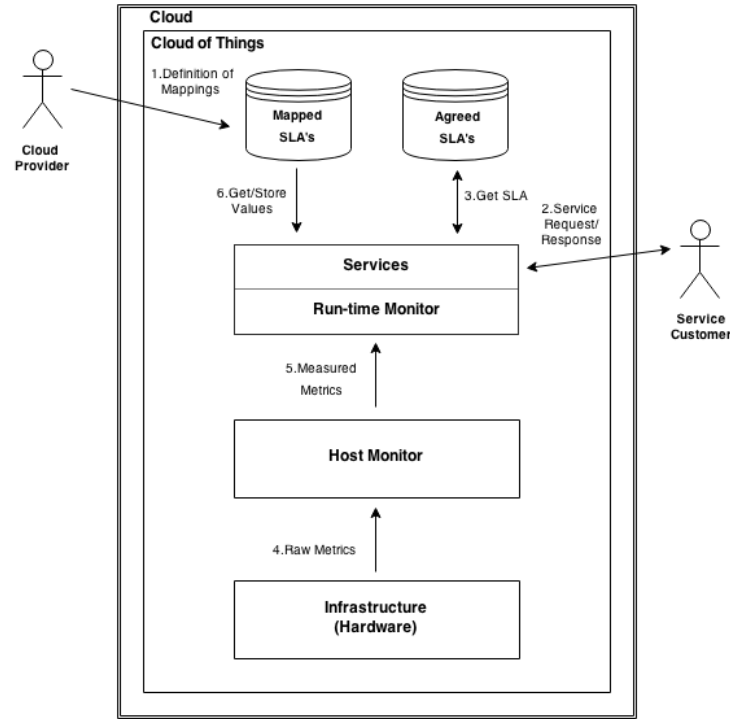


Fig. 5: Cloud of Things Architecture.

The presented architecture is based on the LoM2HiS Framework [26] architecture. In this architecture the *Services* component and the *Run-time Monitor* represents the application layer where services are deployed using a Web Service container. The *Run-time Monitor* is responsible to monitor the services based on the negotiated and agreed SLAs. After the Cloud provider agrees on the SLA terms, the agreed SLAs are stored in the repository for service provisioning and the following steps are executed:

1. The Cloud provider creates rules for the framework mappings using Domain Specific Languages<sup>12</sup> (DSL's).
2. The customer requests the provisioning of an agreed service.
3. Once the request is received, the run-time monitor loads the service SLA from the agreed SLA repository.
4. The resource metrics are measured by monitoring agents, these metrics are stored in a raw format that later are accessed by the host monitor.
5. The host monitor extracts metric-value pairs from the raw metrics and them transmits them periodically to the run-time monitor.

<sup>12</sup> Domain Specific Languages are small languages that normally are tailored to a specific problem domain.



6. After receive the low-level metrics, the run-time monitor uses predefined mapping rules to map the low-level metrics into an equivalent form of the agreed SLA and then the resulting map is stored in the mapped metrics repository.

In this architecture, the *Run-time monitor* uses the mapped values to monitor the status of the deployed services. Once it detects that a SLA is violated, the *Run-time monitor* must alert the customer of the violated SLA. At this point the customer is responsible to take the decisions in order to correct the state of the system.



## 5 Evaluation Methodology

The evaluation of the solution can be performed according two perspectives, one consists in evaluate the performance of the Cloud and the other consists in evaluate the correctness of the received RFID events.

### 5.1 Cloud Performance Evaluation

An aspect that is important concerns with the performance of the Cloud regarding the amount of data generated by the events. Cloud computing creates an illusion that available computing resources on demand are infinite [27]. However, with the increase of the amount of events, we must measuring these computing resources in order to determine if the system is fulfilling the *QoS* requirements. To perform the evaluation of the behaviour and performance of the Cloud we need to measure some system metrics such:

- *CPU Utilization*: CPU Utilization indicates the percentage of time that the CPU was working at the instances in the Cloud. Normally, this metric is available through the Cloud providers. Usually, the range of this metrics is given in percentage that can vary between 0-100%.
- *Memory Usage*: Memory Usage indicates the amount of memory that is consumed by the system in a given period of time. The range of this metric is given in MBytes.
- *System Load*: System Load is metric that indicates the general state of the system. This metric estimates the general performance of the system by measuring the number of received events. The range of this metric vary between 0 and 1. When this metric has a value of 0, it means that the system is not receiving any events at the time. If this metric has a value of 1, it means that the system is overloaded, consequently the CPU Utilization and Memory Usage metrics are close to the maximum value.

### 5.2 RFID Evaluation

RFID has its particular characteristics, which makes that traditional event processing systems cannot support them. Furthermore, RFID events are temporal constrained [28]. Temporal constraints as the time interval between two events and the time interval for a single event are critical to event detection. In order to assure the correctness of these received events, some conditions must be defined regarding the time interval that between the events. Temporal constraints are not the only cause that can compromise the correctness of received the RFID events such as collisions on the air interface, tag detuning and tag misalignment [29].

The metric used to evaluate the RFID performance is proposed by Correia [30]. The *PresenceRate* is a metric that takes in account the reported time of reading a RFID tag. This metric measures the ratio between the time that a

given object spent in a given area and the expected spent time. Correia et al. [30] defines that the expected behaviour of this metric is given by the following values:

- $PresenceRate = 0$ : the object was not detected by the system.
- $0 < PresenceRate < 1$ : there are false negative readings and they were reported.
- $PresenceRate = 1$ : ideal scenario where the reported time is exactly the same as expected.
- $PresenceRate > 1$ : there are false positive readings over-estimating the time spent by the object in the read area.

In order to perform a general evaluation of the system the we can establish a relation between the *PresenceRate* and the metrics used to evaluate the Cloud performance. The scenario where the *PresenceRate* is larger then 1 can be used to determine if the system is overloaded. In this scenario the metrics of *PresenceRate* and *SystemLoad* must be compared to verify this condition. If the *SystemLoad* value is equal to 1, it means that the system is operating in it maximum capacity and the amount of resources available is not enough to process all the requests. Otherwise this comparison can determine of the system is overpowered. If the *PresenceRate* is equal to 1 and the *SystemLoad* is close to 0, it means that the available resources are more than the required to process all the requests.



## 6 Conclusion

The Internet of Things is a paradigm that will revolutionize the way in that common objects interact with the environment. In this document we introduce the concept of Internet of Things in the scope of smart places. A smart place is an ecosystem that is composed by smart objects, in this case RIFD tagged objects, that are interconnected within the Internet infrastructure present in this smart place. However, due to the diversity of smart objects and its different communication protocols, these smart places are characterized for its heterogeneity which increases the complexity the execution of management tasks such as deployment and monitoring of the IoT application that is running in this smart place. Actually the state of the art solution reduces the complexity of such tasks by automating the deployment process, but as mentioned in 2 this solution requires an high level of expertise in order to perform this tasks.

Therefore, with this work we pretend to decrease the complexity of the execution of these management tasks in two ways. First, we pretend to automate the deployment of IoT applications in smart places by using Cloud orchestration tools, these tools will also allow to describe the application structure in a high-level perspective, thus modelling the applications structure only requires to have in mind the application logic. In the other hand, the main objective of this work is to enable the business owners to define the Service Level Agreements (SLA) only having in mind the business rules of its smart place. With this approach the business owner will be able to monitoring the performance of its smart place and establish the amount of resources needed by the application in order to have an adequate *QoS*.

In the evaluation of the developed solution, we will evaluate the performance of the application regarding the amount of data that is generated by the objects that are in the smart place. An important point of the evaluation is to demonstrate the possible scenarios where the smart place generates more data then the Cloud can process and the opposite scenario where the Cloud is overpowered. Ultimately, we pretend that this work contributes to improve the developed work in the field of smart places, focusing in a area that is essential to guarantee the correct operation of them, *QoS* support.



## A Planning

In this section we propose a scheduling that estimates the required time in future work that will be realized.

Milestone	Dates
Create the model of the solution's structure	10/02/2015 - 01/03/2015
Implementation of the solution	02/03/2015 - 15/04/2015
Perform the tests of the evaluation	16/04/2015 - 26/04/2015
Process the results of the evaluation	27/04/2015 - 10/04/2015
Write the dissertation document	11/04/2015 - 15/06/2015
Deliver the final version of the dissertation document	16/06/2015

Table 1: Milestones and corresponding dates for the work of Master Thesis



## References

1. M. Weiser, R. Gold, and J. S. Brown, "The origins of ubiquitous computing research at parc in the late 1980s," *IBM systems journal*, vol. 38, no. 4, pp. 693–696, 1999.
2. J. Gubbi, R. Buyya, S. Marusic, and M. Palaniswami, "Internet of things (iot): A vision, architectural elements, and future directions," *Future Generation Computer Systems*, vol. 29, no. 7, pp. 1645–1660, 2013.
3. M. Weiser, "The computer for the 21st century," *Scientific american*, vol. 265, no. 3, pp. 94–104, 1991.
4. A. Greenfield, *Everyware: The dawning age of ubiquitous computing*. New Riders, 2010.
5. Y. Rogers, "Moving on from weiser's vision of calm computing: Engaging ubicomp experiences," in *UbiComp 2006: Ubiquitous Computing*, pp. 404–421, Springer, 2006.
6. R. Cáceres and A. Friday, "Ubicomp systems at 20: Progress, opportunities, and challenges," *IEEE Pervasive Computing*, vol. 11, no. 1, pp. 14–21, 2012.
7. F. Mattern and C. Floerkemeier, "From the internet of computers to the internet of things," in *From active data management to event-based systems and more*, pp. 242–259, Springer, 2010.
8. L. Atzori, A. Iera, and G. Morabito, "The internet of things: A survey," *Computer networks*, vol. 54, no. 15, pp. 2787–2805, 2010.
9. D. Cook and S. Das, *Smart environments: technology, protocols and applications*, vol. 43. John Wiley & Sons, 2004.
10. D. Uckelmann, M. Harrison, and F. Michahelles, *Architecting the Internet of Things*. Springer Publishing Company, Incorporated, 1st ed., 2011.
11. R. Buyya, C. S. Yeo, S. Venugopal, J. Broberg, and I. Brandic, "Cloud computing and emerging it platforms: Vision, hype, and reality for delivering computing as the 5th utility," *Future Generation computer systems*, vol. 25, no. 6, pp. 599–616, 2009.
12. W.-T. Tsai, X. Sun, and J. Balasooriya, "Service-oriented cloud computing architecture," in *Information Technology: New Generations (ITNG), 2010 Seventh International Conference on*, pp. 684–689, IEEE, 2010.
13. L. M. Vaquero, L. Roderó-Merino, J. Cáceres, and M. Lindner, "A break in the clouds: towards a cloud definition," *ACM SIGCOMM Computer Communication Review*, vol. 39, no. 1, pp. 50–55, 2008.
14. B. Sotomayor, R. S. Montero, I. M. Llorente, and I. Foster, "Virtual infrastructure management in private and hybrid clouds," *Internet Computing, IEEE*, vol. 13, no. 5, pp. 14–22, 2009.
15. S. Zhang, S. Zhang, X. Chen, and X. Huo, "Cloud computing research and development trend," in *Future Networks, 2010. ICFN'10. Second International Conference on*, pp. 93–97, IEEE, 2010.
16. J. O'Sullivan, D. Edmond, and A. Ter Hofstede, "What's in a service?," *Distributed and Parallel Databases*, vol. 12, no. 2-3, pp. 117–133, 2002.
17. L. Zeng, B. Benatallah, A. H. Ngu, M. Dumas, J. Kalagnanam, and H. Chang, "Qos-aware middleware for web services composition," *Software Engineering, IEEE Transactions on*, vol. 30, no. 5, pp. 311–327, 2004.
18. D. Guinard, C. Floerkemeier, and S. Sarma, "Cloud computing, rest and mashups to simplify rfid application development and deployment," in *Proceedings of the Second International Workshop on Web of Things*, p. 9, ACM, 2011.

19. S. Distefano, G. Merlino, and A. Puliafito, "Enabling the cloud of things," in *Innovative Mobile and Internet Services in Ubiquitous Computing (IMIS), 2012 Sixth International Conference on*, pp. 858–863, IEEE, 2012.
20. J. Zhou, T. Leppanen, E. Harjula, M. Ylianttila, T. Ojala, C. Yu, and H. Jin, "Cloudthings: A common architecture for integrating the internet of things with cloud computing," in *Computer Supported Cooperative Work in Design (CSCWD), 2013 IEEE 17th International Conference on*, pp. 651–657, IEEE, 2013.
21. F. Li, M. Vogler, M. Claeßens, and S. Dustdar, "Towards automated iot application deployment by a cloud-based approach," in *Service-Oriented Computing and Applications (SOCA), 2013 IEEE 6th International Conference on*, pp. 61–68, IEEE, 2013.
22. U. Breitenbücher, T. Binz, K. Képes, O. Kopp, F. Leymann, and J. Wettinger, "Combining declarative and imperative cloud application provisioning based on toasca," *IC2E. IEEE*, 2014.
23. T. Waizenegger, M. Wieland, T. Binz, U. Breitenbücher, F. Haupt, O. Kopp, F. Leymann, B. Mitschang, A. Nowak, and S. Wagner, "Policy4tosca: A policy-aware cloud service provisioning approach to enable secure cloud computing," in *On the Move to Meaningful Internet Systems: OTM 2013 Conferences*, pp. 360–376, Springer, 2013.
24. U. Breitenbücher, T. Binz, C. Fehling, O. Kopp, F. Leymann, and M. Wieland, "Policy-aware provisioning and management of cloud applications," *International Journal On Advances in Security*, vol. 7, no. 1 and 2, pp. 15–36, 2014.
25. U. Breitenbücher, T. Binz, O. Kopp, F. Leymann, and M. Wieland, "Policy-aware provisioning of cloud applications," in *SECURWARE 2013, The Seventh International Conference on Emerging Security Information, Systems and Technologies*, pp. 86–95, 2013.
26. V. C. Emeakaroha, I. Brandic, M. Maurer, and S. Dustdar, "Low level metrics to high level sla-som2his framework: Bridging the gap between monitored metrics and sla parameters in cloud environments," in *High Performance Computing and Simulation (HPCS), 2010 International Conference on*, pp. 48–54, IEEE, 2010.
27. M. Armbrust, O. Fox, R. Griffith, A. D. Joseph, Y. Katz, A. Konwinski, G. Lee, D. Patterson, A. Rabkin, I. Stoica, *et al.*, "M.: Above the clouds: A berkeley view of cloud computing," 2009.
28. F. Wang, S. Liu, P. Liu, and Y. Bai, "Bridging physical and virtual worlds: complex event processing for rfid data streams," in *Advances in Database Technology-EDBT 2006*, pp. 588–607, Springer, 2006.
29. C. Floerkemeier and M. Lampe, "Issues with rfid usage in ubiquitous computing applications," in *Pervasive Computing*, pp. 188–193, Springer, 2004.
30. N. Correia, "RFIDToys: A Flexible Testbed Framework for RFID Systems," Master's thesis, Instituto Superior Tcnico, Portugal, 2014.